

# CS5344

## Clustering High Dimensional Data




# High Dimensional Data

Given a cloud of data points we want to understand its structure




# Problem of Clustering




- Given a set of points, with a notion of distance between points, group the points into some number of **clusters** such that
  - Members of a cluster are similar to each other
  - Members of different clusters are dissimilar
- **Points are in a high-dimensional space**
  - Each dimension corresponds to a feature/attribute
- **Similarity is defined using a distance measure**
  - Euclidean, Cosine, Jaccard, Edit distance, ...

# Example: Clusters & Outliers



# Clustering is a Hard Problem!



# Why is Clustering Hard?

- **Clustering in two dimensions looks easy**
- **Clustering small amounts of data looks easy**
- **BUT**
  - Number of clusters is typically not known
  - Clusters may be of arbitrary shapes and sizes
  - Quality of clustering result
    - Depends on similarity measure and method
    - Measured by its ability to find some or all hidden patterns
- **Many applications operate in high dimensional space (not 2, but 10 or 10000 dimensions)**
  - Almost all pairs of points are at about the same distance!

# Application: Galaxies

- Catalog of 2 billion “sky objects” represents objects by their radiation in 7 dimensions (frequency bands)
- Task: Cluster into similar objects, e.g., galaxies, nearby stars, quasars, etc.
- Sloan Digital Sky Survey



# Application: Music CDs

- Music divided into categories, and customers prefer a few categories. What are categories actually?
- Represent a CD by a set of customers who bought it
  - Similar CDs have similar sets of customers
- **Task: Find clusters of similar CDs**
  - Space of all CDs, one dimension for each customer
  - Values in a dimension may be 0 or 1
  - A CD is a point in this space  $(x_1, x_2, \dots, x_k)$  where  $x_i = 1$  iff the  $i^{\text{th}}$  customer bought the CD
- **For Amazon, the dimension is tens of millions**



# Application: Cluster Documents

- Task: Find topics
- Represent each document as a vector  $(x_1, x_2, \dots, x_k)$  where  $x_i = 1$  iff the word  $v_i$  appears in the document
- Documents with similar sets of words may be about the same topic

# Overview

- **Distance Measures**
- **Clustering Algorithms**
  - Hierarchical or agglomerative
  - Point assignment or *k*-means
- **Scaling Up Clustering Algorithms**
  - BFR
  - CURE

# Distance Measure

- Each clustering problem is based on some notion of distance between objects or points
- **Euclidean Distance**
  - Based on locations of points in a  $d$ -dimensional space
  - Points are vectors of real numbers
  - Length of vector is  $d$
- **Non-Euclidean Distance**
  - Based on the properties of points

# Distance Measure

- **Distance measure is a function  $d(x,y)$  between two points  $x$  and  $y$  which satisfies the following axioms**
  1. Non-negativity:  $d(x, y) \geq 0$
  2. Identity:  $d(x, y) = 0$  if and only if  $x = y$
  3. Symmetry:  $d(x, y) = d(y, x)$
  4. Triangle Inequality:  $d(x, y) \leq d(x, z) + d(z, y)$

# Examples of Distance Measure

- **$L_p$ -norm**

$$d(\mathbf{x}, \mathbf{y}) = \left( \sum_i (x_i - y_i)^p \right)^{\frac{1}{p}}$$

- **$L_2$ -norm = Euclidean Distance**

- Square the distance in each dim, sum the squares and take sqrt
- E.g. two points in n-dimensional space
- $d(\mathbf{x}, \mathbf{y}) = \text{sqrt} \left( (x_1 - y_1)^2 + (x_2 - y_2)^2 + \dots + (x_n - y_n)^2 \right)$

- **$L_1$ -norm = Manhattan Distance**

- Sum of the magnitude of differences in each dimension
- $d(\mathbf{x}, \mathbf{y}) = |x_1 - y_1| + |x_2 - y_2| + \dots + |x_n - y_n|$
- Distance to travel if constrained along grid lines

# Examples of Distance Measure

- **Cosine Distance**
  - Points are vectors with integer or boolean values

$$\text{cosine}(\mathbf{x}, \mathbf{y}) = \frac{\sum_i x_i y_i}{\sqrt{\sum_i x_i^2} \sqrt{\sum_i y_i^2}}$$

# Examples of Distance Measure

- **Jaccard Distance**

- If A and B are two sets, then



$$J(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

- **Edit Distance**

- Points are strings
  - Minimum number of inserts and deletes of characters needed to convert one string into another
  - E.g. if  $x = abcde$  and  $y = acfdeg$ ,  $d(x, y) = 3$

# Clustering Strategies



- **Hierarchical or Agglomerative**
  - Each point is a cluster
  - Repeatedly combine the two “nearest” clusters into one
- **Point Assignment**
  - Maintain a set of clusters
  - Points belong to “nearest” cluster






# Hierarchical Clustering

- **Key operation: Repeatedly combine two nearest clusters**
- **Produce a set of nested clusters as a hierarchical tree**
- **Visualized as a dendrogram (tree-like diagram records the sequence of merges)**




# Hierarchical Clustering

- Does not assume any number of clusters
- Any desired number of clusters can be obtained by 'cutting' the dendrogram at the proper level
- May correspond to meaningful taxonomies
  - e.g., animal kingdom, phylogeny reconstruction in biological sciences



# Hierarchical Clustering

- **Key operation: Repeatedly combine two nearest clusters**
- **Three important questions:**
  1. How to represent a cluster of more than one point?
  2. How to determine the “nearness” of clusters?
  3. When to stop combining clusters?





# How to Represent Cluster?


- Each point is a cluster initially
- As we merge clusters, how to represent the “location” of each cluster?
  - Need this information to know which pair of clusters is closest
- **Euclidean space**
  - Each cluster has a **centroid** = average of its (data) points

# Which 2 Clusters to Merge?

- **Euclidean space**
  - Measure cluster distances by distances of centroids



$\sigma$  ... data point  
 $\bar{x}$  ... centroid




# Non-Euclidean Space

- There is no “average” of two points
- The only “locations” are the points themselves
- **Clustroid** = (data) point “closest” to other points
  - Treat clustroid as if it were centroid, when computing inter-cluster distances
- **Possible meanings of “closest”**
  - Smallest maximum distance to other points
  - Smallest average distance to other points
  - Smallest sum of squares of distances to other points

# Non-Euclidean Space

- For distance metric  $d$ , clustroid  $c$  of cluster  $C$  is

$$\min_c \sum_{x \in C} d(x, c)^2$$




**Cluster on 3 datapoints**

**Centroid** is the avg. of all (data)points in the cluster. This means centroid is an “artificial” point.

**Clustroid** is an **existing** (data)point that is “closest” to all other points in the cluster.

# “Nearness” of Clusters *Different measures*

1. Based on distance between clustroids/centroids
2. Intercluster distance = minimum of the distances between any two points, one from each cluster
3. Pick a notion of “cohesion” of clusters and merge cluster whose union is most cohesive
  - **Diameter** of merged cluster = maximum distance between points in the cluster
  - **Average distance** between points in the cluster
  - **Density-based** (take diameter or avg. distance, and divide by number of points in cluster)





# When to Terminate?

- Pre-determined *number of clusters*
- When merging two clusters leads to a “bad” cluster
  - Diameter of merged cluster exceeds some threshold
  - Diameter exceeds average diameter by a wide margin
  - Density of cluster falls below some threshold

# Implementation

- **Naïve implementation of hierarchical clustering**
  - At each step, compute pairwise distances between all pairs of clusters, then merge
  - $O(n^3)$
- **Careful implementation using Priority Queue can reduce time to  $O(n^2 \log n)$** 
  - Still too expensive for really big datasets that do not fit in memory

# K-Means Algorithm

- Assumes Euclidean space/distance
- Assumes number of clusters  $k$  is given
- Initialize clusters by picking one point per cluster
  - Pick one point at random, then  $k-1$  other points, each **as far away as possible** from the previous points
  - Make these points the centroid of their clusters

# K-Means Algorithm


1. For each point, place it in the cluster whose current centroid it is nearest
2. After all points are assigned, update the locations of centroids of the  $k$  clusters
3. Reassign all points to their closest centroid
  - Sometimes moves points between clusters

**Repeat 2 and 3 until convergence**

- **Convergence:** Points don't move between clusters and centroids stabilize

# Example


## Clusters after Round 1



x ... data point  
□ ... centroid

# Example


## Clusters after Round 2



x ... data point  
□ ... centroid

# Example


Clusters at the end



x ... data point  
□ ... centroid

# How to Select $k$ ?


- Try different values of  $k$
- Look at the change in the average distance to centroid as  $k$  increases
- Average falls rapidly until right  $k$ , then changes little






# Example


Too few;  
many long  
distances  
to centroid.



Too many;  
little improvement  
in average  
distance.



Just right;  
distances  
rather short.



# Does k-means Converge?



- Most of the time, convergence happens in the first few iterations
- Stopping criterion is often changed to “**until relatively few points change clusters**”
- Complexity is  $O(n*k*m)$ 
  - $n$  = number of points
  - $k$  = number of clusters
  - $m$  = number of iterations

# **BFR Algorithm** **(Bradley-Fayyad-Reina)**

**Extension of *k*-means to large data**

# BFR Algorithm

- Variant of *k*-means designed to handle very large (disk-resident) data sets in high dimensions
- Assumes points in clusters are normally distributed around a centroid in a Euclidean space
  - Standard deviations in different dimensions may vary
  - Clusters are axis-aligned ellipses
- **Efficient way to summarize clusters**
  - Memory required  $O(\text{clusters})$  and not  $O(\text{data})$



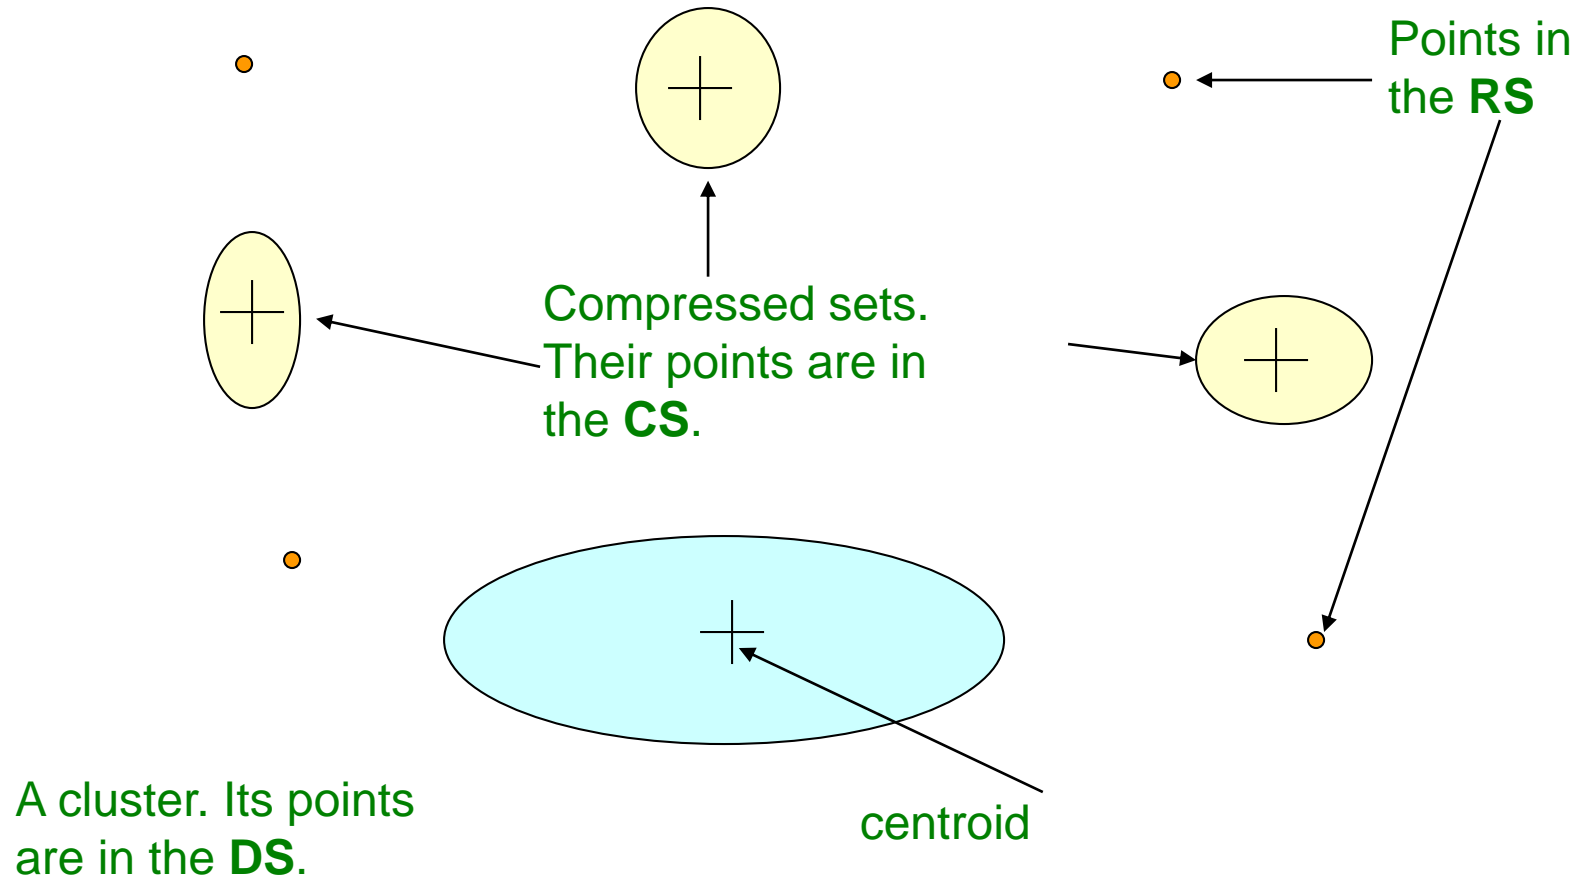
# BFR Algorithm

- Points are read from disk to memory in chunks
- Points from previous memory loads are summarized by simple statistics
- Select initial  $k$  centroids from the first chunk
  - Take  $k$  random points; or
  - Take a small random sample and cluster optimally; or
  - Take a sample; pick a random point, and then  $k-1$  more points, each as far from the previously selected points as possible

# BFR Algorithm

- **Keep track of 3 sets of points in memory**
  - **Discard set (DS):**
    - Points close enough to a centroid to be summarized
  - **Compression set (CS):**
    - Groups of points that are close together but not close to any existing centroid
    - These points are summarized, but not assigned to a cluster
  - **Retained set (RS):**
    - Isolated points to be assigned to a compression set

# BFR: “Galaxies” Picture




**Discard set (DS):** Close enough to a centroid to be summarized

**Compression set (CS):** Summarized, but not assigned to a cluster

**Retained set (RS):** Isolated points

# Summarizing Points


- For each cluster, discard set (DS) is summarized by
  - Number of points,  $N$
  - Vector  $SUM$  whose  $i^{th}$  component is the sum of the coordinates of the points in the  $i^{th}$  dimension
  - Vector  $SUMSQ$  where  $i^{th}$  component = sum of squares of coordinates in  $i^{th}$  dimension





# Summarizing Points

- $2d + 1$  values represent any size cluster
  - $d$  = number of dimensions
- Centroid (Average in each dimension) is given by  $SUM_i / N$ 
  - $SUM_i = i^{th}$  component of SUM
- Variance of a cluster's DS in dimension  $i$  is given by  $(SUMSQ_i / N) - (SUM_i / N)^2$ 
  - Standard deviation is the square root of variance
- **Next step: Actual clustering**



# Processing “Memory-Load” of Points

- Find those points that are “sufficiently close” to a cluster centroid and add those points to that cluster and the DS
  - These points are so close to the centroid that they can be summarized and then discarded
- Use any main-memory clustering algorithm to cluster the remaining points and the old RS
  - Clusters go to the **CS**; outlying points to the **RS**

**Discard set (DS):** Close enough to a centroid to be summarized.

**Compression set (CS):** Summarized, but not assigned to a cluster

**Retained set (RS):** Isolated points

# Processing “Memory-Load” of Points

- **DS set:** Adjust statistics of the clusters to account for the new points
  - Add  $N_s$ ,  $SUM_s$ ,  $SUMSQ_s$
- **Consider merging compressed sets in the CS**
  - Add corresponding values of  $N_s$ ,  $SUM_s$ ,  $SUMSQ_s$
- **Last round, merge all compressed sets in the CS and all RS points into their nearest cluster**

**Discard set (DS):** Close enough to a centroid to be summarized.

**Compression set (CS):** Summarized, but not assigned to a cluster

**Retained set (RS):** Isolated points

# BFR Algorithm

## Two Questions:

1. How do we decide if a point is “close enough” to a cluster that we will add the point to that cluster?
2. How do we decide whether two compressed sets (CS) deserve to be combined into one?

# Mahalanobis Distance

- **Measure of the distance between a point P and a distribution D**
  - How many standard deviations is P away from mean of D?
- **For a point  $(x_1, \dots, x_n)$  and centroid  $(c_1, \dots, c_n)$**


- $$d(x, c) = \sqrt{\sum_{i=1}^n \left( \frac{x_i - c_i}{\sigma_i} \right)^2}$$

*Normalized Euclidean Distance  
from Centroid*

- Normalize in each dimension:  $y_i = (x_i - c_i) / \sigma_i$
- Take sum of the squares of the  $y_i$
- Take the square root
- $\sigma_i$  is the standard deviation of points in the cluster in the  $i^{\text{th}}$  dimension


# Mahalanobis Distance

- Compute MD between point and each cluster centroid
- Choose cluster whose centroid has the least MD
- Add point to cluster if  $MD < \text{threshold}$ 
  - e.g. threshold = 2 (standard deviation)



# Combining Two CS clusters

- Compute the variance of combined cluster
- $N$ ,  $SUM$ , and  $SUMSQ$  allow us to make that calculation quickly
- Combine if the combined variance is below some threshold





# CURE Algorithm

**Extension of  $k$ -means to  
clusters of arbitrary shapes**




# CURE Algorithm

- BFR and *k*-means assume clusters are normally distributed in each dimension
  - Axes are fixed
  - Ellipses at an angle are *not OK*
- **CURE (Clustering Using Representatives)**
  - Assumes Euclidean distance
  - Allows clusters of any shape
  - **Uses a collection of representative points to represent clusters**



# Example Stanford Salaries




# CURE Algorithm


## Pass 1

- **Pick a random sample of points that fit in main memory**
- **Cluster these points hierarchically – group nearest points/clusters**
- **Pick representative points**
  - For each cluster, pick a sample of points, as dispersed as possible
  - From the sample, pick representatives by moving them (say) 20% toward the centroid of the cluster


# Example



# Example




# Example



# CURE Algorithm

## Pass 2

- Rescan the whole dataset and visit each point  $p$  in the data set
- Place it in the “closest cluster”
  - Find the closest representative to  $p$  and assign it to representative's cluster



# Summary

- Given a set of points, with a notion of distance between points, group the points into some number of *clusters*
- Centroid in Euclidean space and clustroid in non-Euclidean space
- Agglomerative hierarchical clustering
- $k$ -means, BFR ( $k$ -means extended for large data sets), CURE ( $k$ -means extended for arbitrary clusters)
- Self-Study: Clustering in non-Euclidean space
  - GRGPF